



Vivekanand Education Society's Institute of Technology

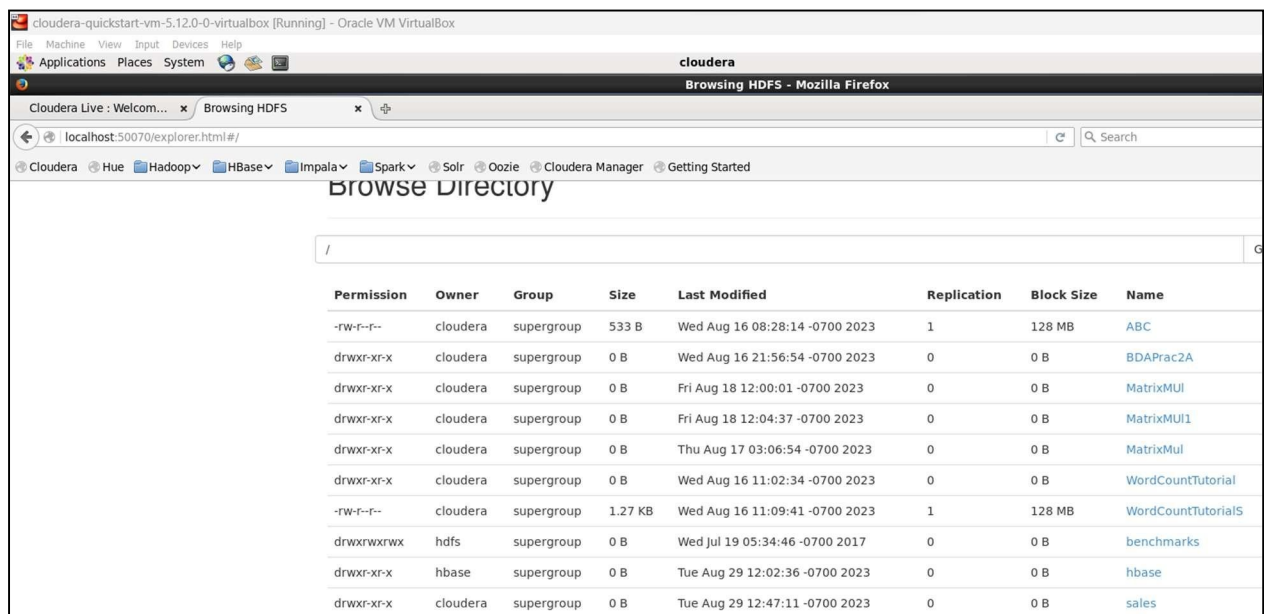
Approved by AICTE & Affiliated to University of Mumbai

Artificial Intelligence and Data Science Department

Big Data Analytics/Odd Sem 2023-23/Experiment 3

Importing tables from RDMS to HDFS using Sqoop:

```
[cloudera@quickstart ~]$ sqoop import --connect jdbc:mysql://localhost/sales --username=root --password="cloudera" --table=sales1 --target-dir=/sales/sales --incremental append --check-column month_number --fields-terminated-by='\t';
```

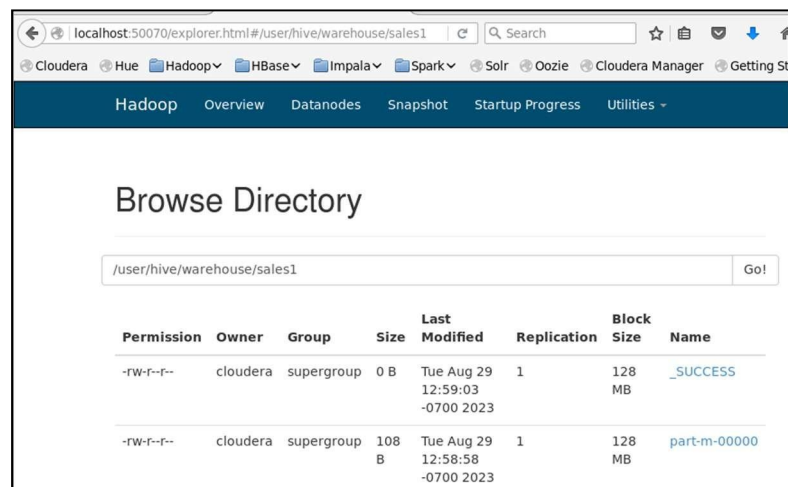


The screenshot shows the Cloudera Live interface with a web browser displaying the HDFS directory listing. The address bar shows 'localhost:50070/explorer.html#/'. The page title is 'Browse Directory'. The table below lists the contents of the HDFS directory.

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	cloudera	supergroup	533 B	Wed Aug 16 08:28:14 -0700 2023	1	128 MB	ABC
drwxr-xr-x	cloudera	supergroup	0 B	Wed Aug 16 21:56:54 -0700 2023	0	0 B	BDAPrac2A
drwxr-xr-x	cloudera	supergroup	0 B	Fri Aug 18 12:00:01 -0700 2023	0	0 B	MatrixMUI
drwxr-xr-x	cloudera	supergroup	0 B	Fri Aug 18 12:04:37 -0700 2023	0	0 B	MatrixMUI1
drwxr-xr-x	cloudera	supergroup	0 B	Thu Aug 17 03:06:54 -0700 2023	0	0 B	MatrixMul
drwxr-xr-x	cloudera	supergroup	0 B	Wed Aug 16 11:02:34 -0700 2023	0	0 B	WordCountTutorial
-rw-r--r--	cloudera	supergroup	1.27 KB	Wed Aug 16 11:09:41 -0700 2023	1	128 MB	WordCountTutorialS
drwxrwxrwx	hdfs	supergroup	0 B	Wed Jul 19 05:34:46 -0700 2017	0	0 B	benchmarks
drwxr-xr-x	hbase	supergroup	0 B	Tue Aug 29 12:02:36 -0700 2023	0	0 B	hbase
drwxr-xr-x	cloudera	supergroup	0 B	Tue Aug 29 12:47:11 -0700 2023	0	0 B	sales

Importing Table From HDFS to HIVE:

```
[cloudera@quickstart ~]$ sqoop import-all-tables --connect jdbc:mysql://localhost/sales --username=root --password "cloudera" --warehouse-dir /user/hive/warehouse
```



The screenshot shows the Cloudera Live interface with a web browser displaying the HDFS directory listing. The address bar shows 'localhost:50070/explorer.html#/user/hive/warehouse/sales1'. The page title is 'Browse Directory'. The table below lists the contents of the HDFS directory.

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	cloudera	supergroup	0 B	Tue Aug 29 12:59:03 -0700 2023	1	128 MB	_SUCCESS
-rw-r--r--	cloudera	supergroup	108 B	Tue Aug 29 12:58:58 -0700 2023	1	128 MB	part-m-00000



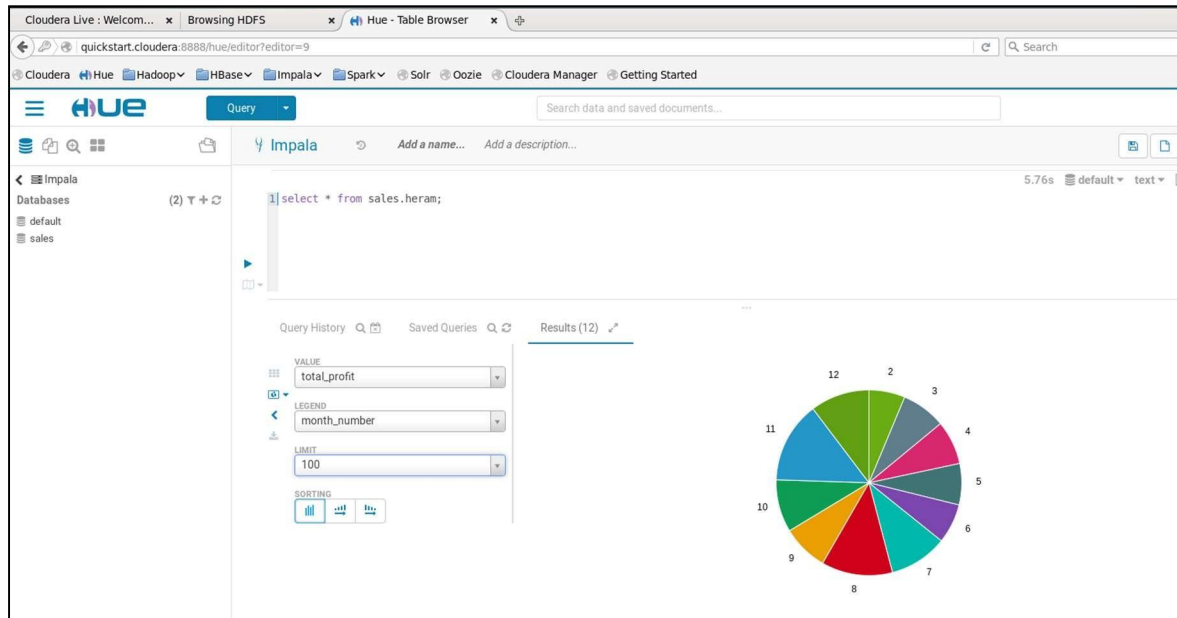
Vivekanand Education Society's Institute of Technology

Approved by AICTE & Affiliated to University of Mumbai

Artificial Intelligence and Data Science Department

Big Data Analytics/Odd Sem 2023-23/Experiment 3

Going to Hue Editor, Importing table, Writing Query And Doing Visualization.



Running Some Queries:

The screenshot shows the Hue Editor interface with a query editor on the left and a table visualization area on the right. The query is:

```
1 SELECT *
2 FROM sales.heram
3 WHERE total_profit>=300200;
```

The table visualization area displays a table with 9 columns: month_number, facecream, facewash, toothpaste, bathingssoap, shampoo, moisturizer, total_units, and tot. The table has 3 rows of data.

	month_number	facecream	facewash	toothpaste	bathingssoap	shampoo	moisturizer	total_units	tot
1	8	3700	1400	5860	9960	2860	1400	36140	361
2	11	2340	2100	7300	13300	2400	2100	41280	412
3	12	2900	1760	7400	14400	1800	1760	30020	300

The screenshot shows the Hue Editor interface with a query editor on the left and a success message on the right. The query is:

```
5
6 Insert into sales.heram values(13,121,345,56,435,43,43,500,234235);
```

The success message is:

✓ Success.



Vivekanand Education Society's Institute of Technology

Approved by AICTE & Affiliated to University of Mumbai

Artificial Intelligence and Data Science Department

Big Data Analytics/Odd Sem 2023-23/Experiment 3

Output:

```
[cloudera@quickstart ~]$ mysql -uroot -pcloudera
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 22
Server version: 5.1.73 Source distribution
```

Copyright (c) 2000, 2013, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its affiliates. Other names may be trademarks of their respective owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

```
mysql> CREATE DATABASE sales;
Query OK, 1 row affected (0.00 sec)
```

```
mysql> use sales;
Database changed
```

```
mysql> LOAD DATA LOCAL Infile '/home/cloudera/Desktop/Heramb/Heram.csv' into table sales1 Fields Terminated By ',' Lines Terminated By '\n';
Query OK, 13 rows affected, 9 warnings (0.02 sec)
Records: 13 Deleted: 0 Skipped: 0 Warnings: 0
```

```
mysql> select * from sales1;
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| month_number | facecream | facewash | toothpaste | bathingsoap | shampoo | moisturizer | total_units | total_profit |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1 | 2500 | 1500 | 5200 | 9200 | 1200 | 1500 | 21100 | 211000 |
| 2 | 2630 | 1200 | 5100 | 6100 | 2100 | 1200 | 18330 | 183300 |
| 3 | 2140 | 1340 | 4550 | 9550 | 3550 | 1340 | 22470 | 224700 |
| 4 | 3400 | 1130 | 5870 | 8870 | 1870 | 1130 | 22270 | 222700 |
| 5 | 3600 | 1740 | 4560 | 7760 | 1560 | 1740 | 20960 | 209600 |
```

```
mysql> show tables
-> ;
+-----+
| Tables_in_sales |
+-----+
| sales1           |
+-----+
```



Vivekanand Education Society's Institute of Technology

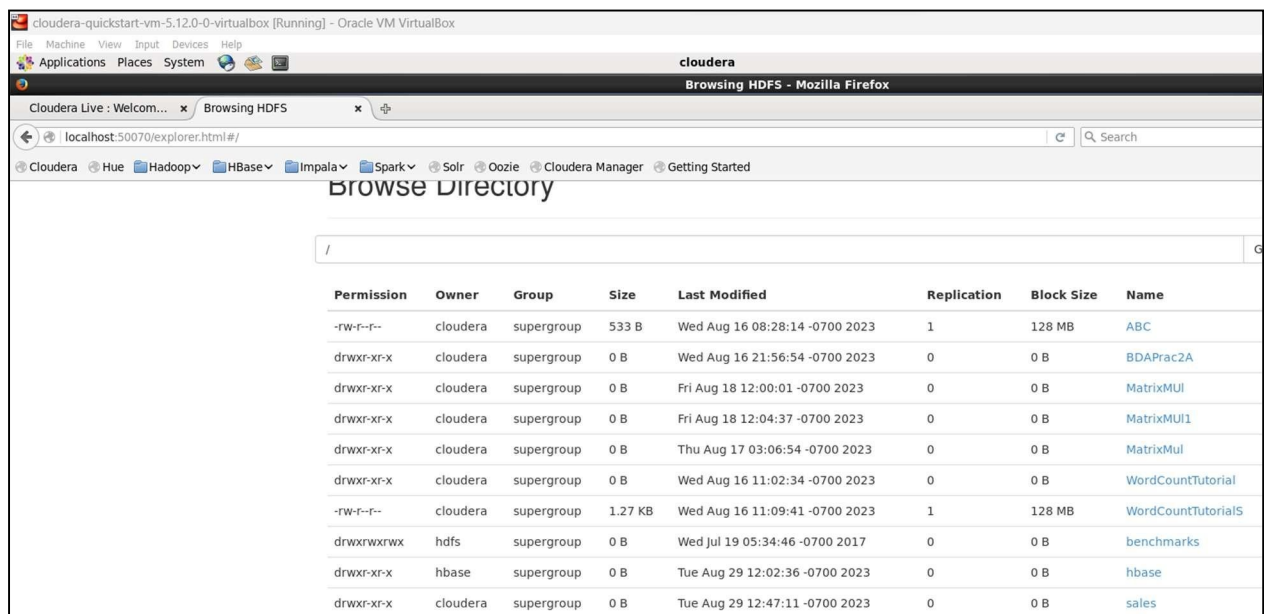
Approved by AICTE & Affiliated to University of Mumbai

Artificial Intelligence and Data Science Department

Big Data Analytics/Odd Sem 2023-23/Experiment 3

Importing tables from RDMS to HDFS using Sqoop:

```
[cloudera@quickstart ~]$ sqoop import --connect jdbc:mysql://localhost/sales --username=root --password="cloudera" --table=sales1 --target-dir=/sales/sales --incremental append --check-column month_number --fields-terminated-by='\t';
```

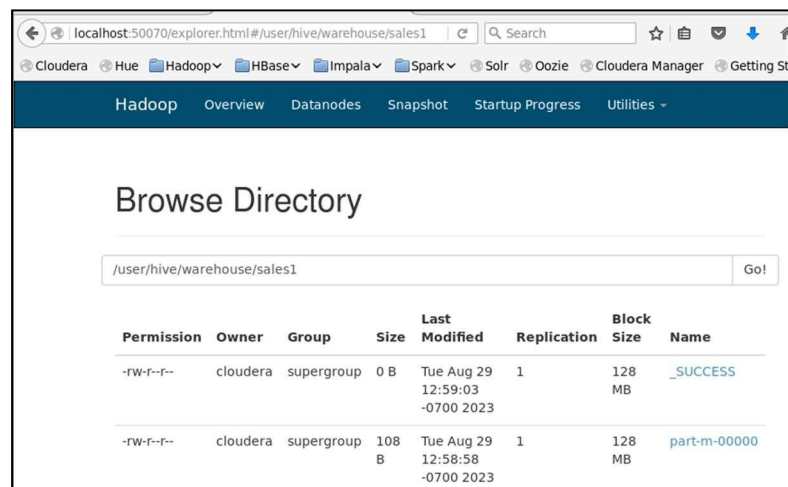


The screenshot shows the Cloudera Live interface with a web browser displaying the HDFS directory listing. The browser address bar shows 'localhost:50070/explorer.html#/'. The page title is 'Browse Directory'. The directory listing shows the following files and directories:

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	cloudera	supergroup	533 B	Wed Aug 16 08:28:14 -0700 2023	1	128 MB	ABC
drwxr-xr-x	cloudera	supergroup	0 B	Wed Aug 16 21:56:54 -0700 2023	0	0 B	BDAPrac2A
drwxr-xr-x	cloudera	supergroup	0 B	Fri Aug 18 12:00:01 -0700 2023	0	0 B	MatrixMUI
drwxr-xr-x	cloudera	supergroup	0 B	Fri Aug 18 12:04:37 -0700 2023	0	0 B	MatrixMUI1
drwxr-xr-x	cloudera	supergroup	0 B	Thu Aug 17 03:06:54 -0700 2023	0	0 B	MatrixMul
drwxr-xr-x	cloudera	supergroup	0 B	Wed Aug 16 11:02:34 -0700 2023	0	0 B	WordCountTutorial
-rw-r--r--	cloudera	supergroup	1.27 KB	Wed Aug 16 11:09:41 -0700 2023	1	128 MB	WordCountTutorialS
drwxrwxrwx	hdfs	supergroup	0 B	Wed Jul 19 05:34:46 -0700 2017	0	0 B	benchmarks
drwxr-xr-x	hbase	supergroup	0 B	Tue Aug 29 12:02:36 -0700 2023	0	0 B	hbase
drwxr-xr-x	cloudera	supergroup	0 B	Tue Aug 29 12:47:11 -0700 2023	0	0 B	sales

Importing Table From HDFS to HIVE:

```
[cloudera@quickstart ~]$ sqoop import-all-tables --connect jdbc:mysql://localhost/sales --username=root --password "cloudera" --warehouse-dir /user/hive/warehouse
```



The screenshot shows the Cloudera Live interface with a web browser displaying the HDFS directory listing for the Hive warehouse. The browser address bar shows 'localhost:50070/explorer.html#/user/hive/warehouse/sales1'. The page title is 'Browse Directory'. The directory listing shows the following files and directories:

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	cloudera	supergroup	0 B	Tue Aug 29 12:59:03 -0700 2023	1	128 MB	_SUCCESS
-rw-r--r--	cloudera	supergroup	108 B	Tue Aug 29 12:58:58 -0700 2023	1	128 MB	part-m-00000



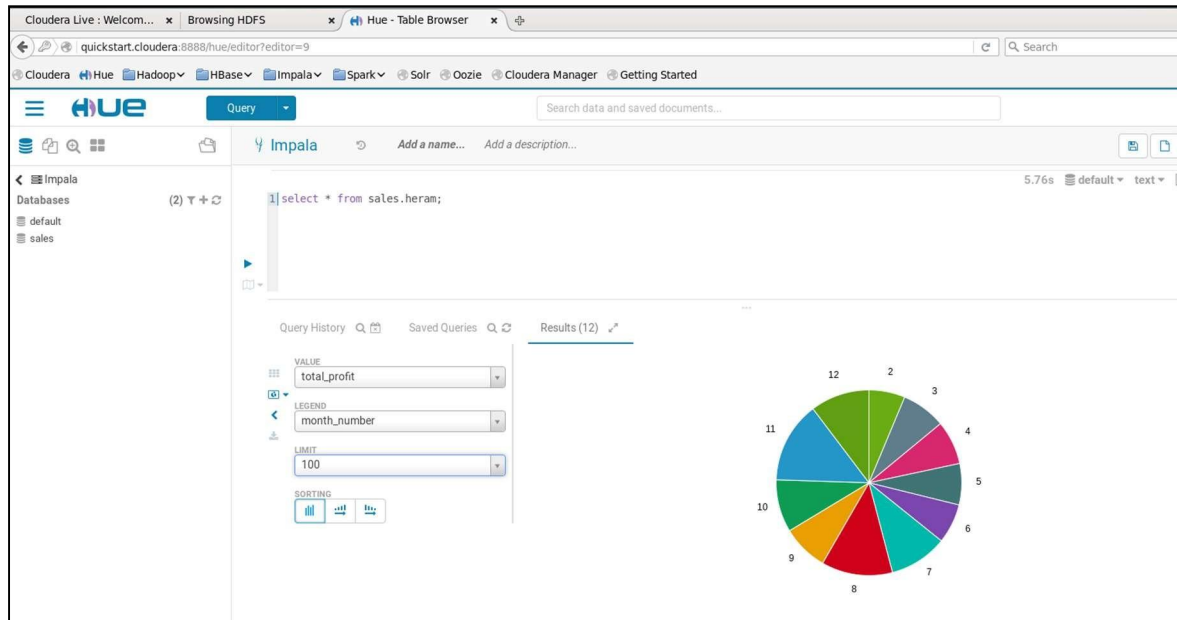
Vivekanand Education Society's Institute of Technology

Approved by AICTE & Affiliated to University of Mumbai

Artificial Intelligence and Data Science Department

Big Data Analytics/Odd Sem 2023-23/Experiment 3

Going to Hue Editor, Importing table, Writing Query And Doing Visualization.



Running Some Queries:

The screenshot shows the Hue Editor interface with a query editor on the left and a table visualization area on the right. The query is:

```
1 SELECT *
2 FROM sales.heram
3 WHERE total_profit>=300200;
```

The table visualization area shows a table with 9 columns: month_number, facecream, facewash, toothpaste, bathingssoap, shampoo, moisturizer, total_units, and tot. The table has 3 rows of data.

	month_number	facecream	facewash	toothpaste	bathingssoap	shampoo	moisturizer	total_units	tot
1	8	3700	1400	5860	9960	2860	1400	36140	361
2	11	2340	2100	7300	13300	2400	2100	41280	412
3	12	2900	1760	7400	14400	1800	1760	30020	300

The screenshot shows the Hue Editor interface with a query editor on the left and a success message on the right. The query is:

```
5
6 Insert into sales.heram values(13,121,345,56,435,43,43,500,234235);
```

The success message is:

✓ Success.



Program:

First Type 'pyspark' in the terminal then type the below commands.

```
>>> sc.appName
u'PySparkShell'

>>> from pyspark import SparkConf, SparkContext

>>> sc
<pyspark.context.SparkContext object at 0x2918c50>

>>> rdd1=sc.textFile("file:/home/cloudera/RT/data1.txt")

>>> rdd2=rdd1.flatMap(lambda line:line.split())

>>> rdd3=rdd2.filter(lambda word:word.startswith('h'))

>>> rdd4=rdd3.map(lambda word:(word,1))

>>> rdd4.collect
```

Output:

```
>>> sc.appName
u'PySparkShell'
>>> from pyspark import SparkConf, SparkContext
>>> sc
<pyspark.context.SparkContext object at 0x1285c50>
>>> rdd1=sc.textFile("file:/home/cloudera/Desktop/BDAPrac2A/Heramb.txt")
>>> rdd2=rdd1.flatMap(lambda line:line.split())
>>> rdd3=rdd2.filter(lambda word:word.startswith('H'))
>>> rdd4=rdd3.map(lambda word:(word,1))
>>> rdd4.collect()
[(u'Hi', 1), (u'Heramb's", 1), (u'Himanshu', 1), (u'Help', 1), (u'He', 1), (u'Help', 1), (u'He', 1), (u'Help', 1)]
```

```
>>> sc.appName
u'PySparkShell'
>>> from pyspark import SparkConf, SparkContext
>>> sc
<pyspark.context.SparkContext object at 0x1285c50>
>>> rdd1=sc.textFile("file:/home/cloudera/Desktop/BDAPrac2A/Heramb.txt")
>>> rdd2=rdd1.flatMap(lambda line:line.split())
>>> rdd3=rdd2.filter(lambda word:word.startswith('A'))
>>> rdd4=rdd3.map(lambda word:(word,1))
>>> rdd4.collect()
[(u'Anjali', 1), (u'Arnav', 1)]
```



Program + Output RDD Programs

A. Selection

```
from pyspark.sql import SQLContext
sqlContext = SQLContext(sc)
df = sqlContext.read.json("/user/cloudera/iris.json")
df.show()
df.select("species").show()
df.select(df['petalLength'], df['species'] + 1).show()
```

petalLength	(species + 1)
null	null
1.4	null
1.4	null
1.3	null
1.5	null
1.4	null
1.7	null
1.4	null
1.5	null

B. Projection

```
>>> from pyspark import SparkContext
>>> c=sc.parallelize([["name","gender","age"],["A","Male","20"],["B","Female","21"],["C","Male","23"],["D","Female","25"]])
>>> c.collect()
[['name', 'gender', 'age'], ['A', 'Male', '20'], ['B', 'Female', '21'], ['C', 'Male', '23'], ['D', 'Female', '25']]
>>> test=c.map(lambda x: x[0])
>>> print "projection ->%s" %(test.collect())
projection ->['name', 'A', 'B', 'C', 'D']
>>> test=c.map(lambda x:x[1])
>>> print "projection ->%s" %(test.collect())
projection ->['gender', 'Male', 'Female', 'Male', 'Female']
```




C. Union

```
>>> sqlContext=SQLContext(sc)
>>> valuesB=[('abc',1),('pqr',2),('mno',7),('xyz',9)]
>>> TableB=sqlContext.createDataFrame(valuesB,['name','customerid'])
>>> valuesC=[('abc',1),('pqr',2),('mno',7),('efg',10),('hik',12)]
>>> TableC=sqlContext.createDataFrame(valuesC,['name','customerid'])
>>> result=TableB.unionAll(TableC)
>>> result.show()
```

name	customerid
abc	1
pqr	2
mno	7
xyz	9
abc	1
pqr	2
mno	7
efg	10
hik	12

D. Aggregate and Grouping

Sum:

```
>>> data=[[1,2],[2,1],[4,3],[4,5],[5,4],[1,4],[1,1]]
>>> list1=sc.parallelize(data)
>>> list1.collect()
[[1, 2], [2, 1], [4, 3], [4, 5], [5, 4], [1, 4], [1, 1]]
>>>
>>>
>>> mapped_list=list1.map(lambda x: (x[0],x[1]))
>>> summation=mapped_list.reduceByKey(lambda x,y: x+y)
>>> summation.collect()
[(1, 7), (2, 1), (4, 8), (5, 4)]
```



Big Data Analytics/Odd Sem 2023-23/Experiment 2B

Average:

```
>>> input1=sqlContext.createDataFrame([(1,2),(2,6),(1,8),(2,4),(3,1),(3,1),(3,1)],["col1","col2"])
>>> input1.groupBy("col1").agg({"col2":"avg"}).show()
+-----+
|col1|avg(col2)|
+-----+
| 1|      5.0|
| 2|      5.0|
| 3|      1.0|
+-----+
```

```
>>> from pyspark.sql import SQLContext
>>> sqlContext = SQLContext(sc)
>>> df = sqlContext.read.json("/user/cloudera/iris.json")
>>> df.groupBy("species").agg({"petalLength": "avg"}).show()
+-----+
| species| avg(petalLength)|
+-----+
|versicolor|      4.26|
| setosa|1.4620000000000002|
| virginica|      5.552|
| null|      null|
+-----+
```

Count:

```
>>> mapped_count = df.map(lambda x : (x[-1],1))
>>> count = mapped_count.reduceByKey(lambda x,y : x+y)
>>> count.collect()
[(None, 2), (u'setosa', 50), (u'versicolor', 50), (u'virginica', 50)]
```



Max & Min element

```
>>> max_element=mapped_list.reduceByKey(lambda x,y:max(x,y))
>>> max_element.collect()
[(1, 4), (2, 1), (4, 5), (5, 4)]
>>>
>>> min_element=mapped_list.reduceByKey(lambda x,y:min(x,y))
>>> min_element.collect()
[(1, 1), (2, 1), (4, 3), (5, 4)]
```

E. Join

```
>>> valueA=[('Pasta',1),('Pizza',2),('Spaghetti',3),('Rice',4)]
>>> rdd1=sc.parallelize(valueA)
>>> TableA=sqlContext.createDataFrame(rdd1,['name','id'])
>>>
>>>
>>> valueB=[('White',1),('Red',2),('Pasta',3),('Spaghetti',4)]
>>> rdd2=sc.parallelize(valueB)
>>> TableB=sqlContext.createDataFrame(rdd2,['name','id'])
>>>
>>> TableA.show()
+-----+
|   name| id|
+-----+
|   Pasta|  1|
|   Pizza|  2|
|Spaghetti|  3|
|    Rice|  4|
+-----+

>>>
>>> TableB.show()
+-----+
|   name| id|
+-----+
|   White|  1|
|    Red|  2|
|   Pasta|  3|
|Spaghetti|  4|
+-----+

>>> ta=TableA.alias('ta')
>>> tb=TableB.alias('tb')
```



```
>>> inner_join=ta.join(tb,ta.name==tb.name)
>>> inner_join.show()
+-----+-----+-----+-----+
|   name| id|   name| id|
+-----+-----+-----+-----+
|Spaghetti| 3|Spaghetti| 4|
|   Pasta| 1|   Pasta| 3|
+-----+-----+-----+-----+

>>>
>>> left=ta.join(tb,ta.name==tb.name,how='left')
>>> left.show()
+-----+-----+-----+-----+
|   name| id|   name| id|
+-----+-----+-----+-----+
|   Rice| 4|   null| null|
|Spaghetti| 3|Spaghetti| 4|
|   Pasta| 1|   Pasta| 3|
|   Pizza| 2|   null| null|
+-----+-----+-----+-----+

>>> right=ta.join(tb,ta.name==tb.name,how='right')
>>> right.show()
+-----+-----+-----+-----+
|   name| id|   name| id|
+-----+-----+-----+-----+
|Spaghetti| 3|Spaghetti| 4|
|   null| null|   White| 1|
|   Pasta| 1|   Pasta| 3|
|   null| null|   Red| 2|
+-----+-----+-----+-----+
```



F. Intersection

```
>>> input1=sc.textFile("file:/home/cloudera/Desktop/BDAPrac2A/input1.txt")
>>> mapinput1=input1.flatMap(lambda x:x.split(","))
>>> mapinput1.collect()
[u'Hello', u' this', u' is', u' Heramb', u' Practical']
>>>
>>> input2=sc.textFile("file:/home/cloudera/Desktop/BDAPrac2A/input2.txt")
>>> mapinput2=input2.flatMap(lambda x:x.split(","))
>>> mapinput2.collect()
[u'Hello', u' this', u' is', u' Heramb', u' Assignment']
>>>
>>>
>>> input3=mapinput1+mapinput2
>>> input3.collect()
[u'Hello', u' this', u' is', u' Heramb', u' Practical', u'Hello', u' this', u' i
s', u' Heramb', u' Assignment']
>>>
>>>
>>> finalintersection=input3.map(lambda word:(word,1))
>>> finalintersection.collect()
[(u'Hello', 1), (u' this', 1), (u' is', 1), (u' Heramb', 1), (u' Practical', 1),
(u'Hello', 1), (u' this', 1), (u' is', 1), (u' Heramb', 1), (u' Assignment', 1)
]
>>> joiningfinalintersection=finalintersection.reduceByKey(lambda x,y:(x+y))
>>> joiningfinalintersection.collect()
[(u' Heramb', 2), (u' Assignment', 1), (u' this', 2), (u' Practical', 1), (u' is
', 2), (u'Hello', 2)]
>>>
>>>
>>> finalans=joiningfinalintersection.filter(lambda x:x[1]>1)
>>> finalans.collect()
[(u'_Heramb', 2), (u' this', 2), (u' is', 2), (u'Hello', 2)]
```




Program:

A. To implement the word Count

```
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable>{

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context
            ) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }
}
```



```
public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
        Context context
        ) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```



Output

```
[cloudera@quickstart BDAPrac2A]$ hadoop dfs -cat /BDAPrac2A/Output/*
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.

After      1
Completed      1
Efforts 1
Hardwork.      1
Heramb's      1
It      1
Lot      1
Of      1
Practical.      1
This      1
Was      1
and      1
is      1
```

B. Matrix-Vector multiplication

Programs:

```
public class MatrixVectorMultiplication {
    public static void main(String[] args) {
        // Define the matrix and vector
        int[][] matrix = {{1, 2, 3}, {4, 5, 6}, {7, 8, 9}};
        int[] vector = {2, 3, 4};

        // Check if matrix and vector dimensions are compatible
        int matrixRows = matrix.length;
        int matrixCols = matrix[0].length;
        int vectorSize = vector.length;

        if (matrixCols != vectorSize) {
            System.out.println("Matrix and vector dimensions are not compatible for
multiplication.");
            return;
        }

        // Perform matrix-vector multiplication
        int[] result = new int[matrixRows];
        for (int i = 0; i < matrixRows; i++) {
```



```
        for (int j = 0; j < matrixCols; j++) {  
            result[i] += matrix[i][j] * vector[j];  
        }  
    }  
  
    // Display the result  
    System.out.println("Result of matrix-vector multiplication:");  
    for (int i = 0; i < matrixRows; i++) {  
        System.out.println(result[i]);  
    }  
}
```

Output:

```
Result of matrix-vector multiplication:  
20  
47  
74
```

Results and Discussions:

MapReduce is a parallel processing model for large scale data:

Word Count:

Result: Efficiently count words occurrences in texts.

Discussions: Map phase splits text, emits . Reduce phase aggregate counts. Scales well for basic tasks.

Matrix-Vector Multiplication:



Output/Procedure:

1. General Commands:

```
hbase(main):001:0> version
1.2.0-cdh5.12.0, rUnknown, Thu Jun 29 04:42:07 PDT 2017

hbase(main):002:0> status
1 active master, 0 backup masters, 1 servers, 0 dead, 2.0000 average load

hbase(main):003:0> whoami
cloudera (auth:SIMPLE)
groups: cloudera, default
```




BDA/Odd Sem 2023-24/Experiment 4

2. Create Table:

```
hbase(main):004:0> create 'customer', 'info', 'orders'
0 row(s) in 2.5610 seconds

=> Hbase::Table - customer
hbase(main):005:0> list
TABLE
customer
1 row(s) in 0.0320 seconds

=> ["customer"]
```

3. Disable and Enable Table:

```
hbase(main):006:0> disable 'customer'
0 row(s) in 2.4500 seconds

hbase(main):007:0> is_disabled 'customer'
true
0 row(s) in 0.0330 seconds

hbase(main):008:0> enable 'customer'
0 row(s) in 1.2910 seconds

hbase(main):009:0> is_enabled 'customer'
true
0 row(s) in 0.0390 seconds
```

4. Describe Table:

```
hbase(main):010:0> describe 'customer'
Table customer is ENABLED
customer
COLUMN FAMILIES DESCRIPTION
{NAME => 'info', DATA_BLOCK_ENCODING => 'NONE', BLOOMFILTER => 'ROW', REPLICATIO
N_SCOPE => '0', VERSIONS => '1', COMPRESSION => 'NONE', MIN_VERSIONS => '0', TTL
=> 'FOREVER', KEEP_DELETED_CELLS => 'FALSE', BLOCKSIZE => '65536', IN_MEMORY =>
'false', BLOCKCACHE => 'true'}
{NAME => 'orders', DATA_BLOCK_ENCODING => 'NONE', BLOOMFILTER => 'ROW', REPLICAT
ION_SCOPE => '0', VERSIONS => '1', COMPRESSION => 'NONE', MIN_VERSIONS => '0', T
TL => 'FOREVER', KEEP_DELETED_CELLS => 'FALSE', BLOCKSIZE => '65536', IN_MEMORY
=> 'false', BLOCKCACHE => 'true'}
2 row(s) in 0.0620 seconds
```

5. Alter Table (Add a new column family 'feedback'):

```
hbase(main):011:0> alter 'customer', {NAME => 'feedback'}
Updating all regions with the new schema...
1/1 regions updated.
Done.
0 row(s) in 2.0590 seconds
```

6. Insert values in Table:

```
hbase(main):003:0> put 'customer', '1001', 'info:name', 'Shreyas Satre'
0 row(s) in 0.3860 seconds

hbase(main):004:0> put 'customer', '1001', 'info:email', 'shreyas.satre@example.
com'
0 row(s) in 0.0070 seconds

hbase(main):005:0> put 'customer', '1001', 'orders:order1', '2023-09-01'
0 row(s) in 0.0180 seconds
```




BDA/Odd Sem 2023-24/Experiment 4

7. Retrive table data:

```
hbase(main):006:0> get 'customer', '1001'
COLUMN                                CELL
info:email                           timestamp=1695028938326, value=shreyas.satre@example.com
info:name                             timestamp=1695028904288, value=Shreyas Satre
orders:order1                         timestamp=1695028949323, value=2023-09-01
3 row(s) in 0.0440 seconds
```

8. Scan table (List all customers/records):

```
hbase(main):027:0> scan 'customer'
ROW                                  COLUMN+CELL
1001                                column=info:name, timestamp=1695028904288, value=Shreyas Sa
tre
1001                                column=orders:order1, timestamp=1695028949323, value=2023-0
9-01
69                                  column=info:email, timestamp=1695029113986, value=shreyas.s
atre@example.com
69                                  column=info:name, timestamp=1695029132268, value=Shreyas Sa
tre
2 row(s) in 0.1440 seconds
```

9. Aggregate Functions(Count Records):

```
hbase(main):028:0> count 'customer'
2 row(s) in 0.0370 seconds

=> 2
```

10. Delete Cell (Remove email for a customer):

```
hbase(main):026:0> delete 'customer', '1001', 'info:email'
0 row(s) in 0.0380 seconds
```

11. Truncate Table:

```
hbase(main):029:0> truncate 'customer'
Truncating 'customer' table (it may take a while):
- Disabling table...
- Truncating table...
0 row(s) in 3.7880 seconds
```

12. Drop Table:

```
hbase(main):031:0> disable 'customer'
0 row(s) in 2.2900 seconds

hbase(main):032:0> drop 'customer'
0 row(s) in 1.3040 seconds
```




Vivekanand Education Society's Institute of Technology

Approved by AICTE & Affiliated to University of Mumbai

Artificial Intelligence and Data Science Department

Big Data Analytics/Odd Sem 2023-23/Experiment 5

Program & Output:

Download Dataset from : <https://archive.ics.uci.edu/ml/datasets/forest+fires>

Upload Dataset into Cloudera.

```
cloudera@quickstart:~/Desktop/Heramb
File Edit View Search Terminal Help
[cloudera@quickstart ~]$ ls
cloudera-manager Downloads kerberos Pictures Videos
cm_api.py eclipse lib Public workspace
Desktop enterprise-deployment.json Music sales1.java
Documents express-deployment.json parcels Templates
[cloudera@quickstart ~]$ cd desktop
bash: cd: desktop: No such file or directory
[cloudera@quickstart ~]$ cd Desktop
[cloudera@quickstart Desktop]$ cd Heramb
[cloudera@quickstart Heramb]$ hdfs dfs -put forestfires.csv /user/cloudera
```

Opening Hive Shell & Creating ForestFire Table:

```
cloudera@quickstart:~
File Edit View Search Terminal Help
[cloudera@quickstart ~]$ sudo hive

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.p
roperties
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.
hive>
>
>
>
>
> Create External Table forestfire(X INT,Y INT,Month STRING,Day STRING,FFMC
FLOAT,DMC FLOAT,Dc FLOAT,ISI FLOAT, Temp FLOAT, RH INT, Wind FLOAT, Rain FLOAT,A
REA FLOAT)
> Row FORMAT DELIMITED
> FIELDS TERMINATED BY ',';
OK
Time taken: 3.799 seconds
```



Vivekanand Education Society's Institute of Technology

Approved by AICTE & Affiliated to University of Mumbai

Artificial Intelligence and Data Science Department

Big Data Analytics/Odd Sem 2023-23/Experiment 5

Loading Data From Dataset Into ForestFire Table:

```
> LOAD DATA INPATH '/user/cloudera/forestfires.csv' OVERWRITE INTO TABLE forestfire;
Loading data to table default.forestfire
chgrp: changing ownership of 'hdfs://quickstart.cloudera:8020/user/hive/warehouse/forestfire/forestfires.csv': Permission denied. user=root is not the owner of inode=forestfires.csv
chmod: changing permissions of 'hdfs://quickstart.cloudera:8020/user/hive/warehouse/forestfire/forestfires.csv': Permission denied. user=root is not the owner of inode=forestfires.csv
Table default.forestfire stats: [numFiles=1, numRows=0, totalSize=25478, rawDataSize=0]
OK
Time taken: 0.537 seconds
..
```

Executing Queries:

Query 1 : select * from forestfire limit 10;

```
> select * from forestfire limit 10;
OK
NULL    NULL    month   day      NULL    NULL    NULL    NULL    NULL    NULL    N
NULL    NULL    NULL
7        5       mar     fri      86.2    26.2    94.3    5.1     8.2     51      6
.7       0.0     0.0
7        4       oct     tue      90.6    35.4    669.1    6.7     18.0    33      0
.9       0.0     0.0
7        4       oct     sat      90.6    43.7    686.9    6.7     14.6    33      1
.3       0.0     0.0
8        6       mar     fri      91.7    33.3    77.5     9.0     8.3     97      4
.0       0.2     0.0
8        6       mar     sun      89.3    51.3    102.2    9.6     11.4    99      1
.8       0.0     0.0
8        6       aug     sun      92.3    85.3    488.0    14.7    22.2    29      5
.4       0.0     0.0
8        6       aug     mon      92.3    88.9    495.6    8.5     24.1    27      3
.1       0.0     0.0
8        6       aug     mon      91.5    145.4    608.2    10.7    8.0     86      2
.2       0.0     0.0
8        6       sep     tue      91.0    129.5    692.6    7.0     13.1    63      5
.4       0.0     0.0
Time taken: 0.294 seconds, Fetched: 10 row(s)
```



Query 2: : select * from forestfire where x=7 and y=4 limit 10;

```
> select * from forestfire where X=7 and Y=4 limit 10;
```


OK												
7	4	oct	tue	90.6	35.4	669.1	6.7	18.0	33	0.9	0.0	0.0
7	4	oct	sat	90.6	43.7	686.9	6.7	14.6	33	1.3	0.0	0.0
7	4	jun	sun	94.3	96.3	200.0	56.1	21.0	44	4.5	0.0	0.0
7	4	aug	sat	90.2	110.9	537.4	6.2	19.5	43	5.8	0.0	0.0
7	4	aug	sat	93.5	139.4	594.2	20.3	23.7	32	5.8	0.0	0.0
7	4	aug	sun	91.4	142.4	601.4	10.6	16.3	60	5.4	0.0	0.0
7	4	sep	fri	92.4	117.9	668.0	12.2	19.0	34	5.8	0.0	0.0
7	4	sep	mon	90.9	126.5	686.5	7.0	19.4	48	1.3	0.0	0.0
7	4	oct	fri	90.0	41.5	682.6	8.7	11.3	60	5.4	0.0	0.0
7	4	aug	sun	94.8	108.3	647.1	17.0	16.4	47	1.3	0.0	1.56

Time taken: 0.2 seconds, Fetched: 10 row(s)

Query 3: select MONTH, avg(FFMC) as Average from forestfire group by MONTH;

apr	85.7888895670573
aug	92.33695594124173
dec	84.96666717529297
feb	82.90499916076661
jan	50.39999961853027
jul	91.32812428474426
jun	89.42941194422104
mar	89.44444345544886
may	87.3499984741211
month	NULL
nov	79.5
oct	90.45333251953124
sep	91.24302336227062

Time taken: 29.623 seconds, Fetched: 13 row(s)





Vivekanand Education Society's Institute of Technology

Approved by AICTE & Affiliated to University of Mumbai

Artificial Intelligence and Data Science Department

Big Data Analytics/Odd Sem 2023-23/Experiment 5


Query 4: SELECT MONTH , MAX(RH) AS MAXIMUM FROM forestfire GROUP BY MONTH HAVING MONTH ='sep';

```
OK
sep      86
Time taken: 26.654 seconds, Fetched: 1 row(s)
```

Query 5: select DAY, SUM(AREA) AS AREA from forestfire group by DAY ORDER BY DAY;

```
day      NULL
fri      447.24000039696693
mon      706.5299995839596
sat      2144.8599796295166
sun      959.9299972057343
thu      997.1000298261642
tue      807.79000864923
wed      578.5999903082848
Time taken: 45.033 seconds, Fetched: 8 row(s)
```

Chart type X-Axis day Y-Axis area



Query 6: SELECT MONTH, MAX(DC) AS MAXIMUM FROM forestfire GROUP BY MONTH ORDER BY MONTH;

```
apr      97.1
aug      819.1
dec      354.6
feb      353.5
jan      171.4
jul      795.9
jun      433.3
mar      103.8
may      113.8
month    NULL
nov      106.7
oct      696.1
sep      860.6
Time taken: 50.182 seconds, Fetched: 13 row(s)
```



BDA/Odd Sem 2023-23/Experiment 6

```
[cloudera@quickstart ~]$ pyspark
```

```
>>> df = sqlContext.createDataFrame([[0,33.3,-17.5],[1,40.4,-20.5],[2,28.6,-23.9],[3,29.5,-19.0],[4,32.8,-18.84]],["other","lat","long"])
23/10/04 07:12:13 WARN ShortCircuitDomainSocketFactory: The short-circuit local reads feature cannot be used because libhadoop cannot be loaded.
```

```
>>> df.show()
+-----+
|other| lat| long|
+-----+
| 0|33.3| -17.5|
| 1|40.4| -20.5|
| 2|28.6| -23.9|
| 3|29.5| -19.0|
| 4|32.8| -18.84|
+-----+
```

```
>>> from pyspark.ml.feature import VectorAssembler
>>> vecAssembler = VectorAssembler(inputCols = ["lat", "long"], outputCol = "features")
>>> new_df = vecAssembler.transform(df)
>>> new_df.show()
```

```
+-----+-----+-----+
|other| lat| long| features|
+-----+-----+-----+
| 0|33.3| -17.5| [33.3, -17.5]|
| 1|40.4| -20.5| [40.4, -20.5]|
| 2|28.6| -23.9| [28.6, -23.9]|
| 3|29.5| -19.0| [29.5, -19.0]|
| 4|32.8| -18.84| [32.8, -18.84]|
+-----+-----+-----+
```

```
>>> from pyspark.ml.clustering import KMeans
>>> kmeans = KMeans(k=2, seed=1)
>>> model = kmeans.fit(new_df.select('features'))
```

```
>>> from pyspark.ml.clustering import KMeans
>>> kmeans = KMeans(k=2, seed=1)
>>> model = kmeans.fit(new_df.select('features'))
```

```
>>> transformed = model.transform(new_df)
>>> transformed.show()
+-----+-----+-----+-----+
|other| lat| long| features|prediction|
+-----+-----+-----+-----+
| 0|33.3| -17.5| [33.3, -17.5]| 0|
| 1|40.4| -20.5| [40.4, -20.5]| 1|
| 2|28.6| -23.9| [28.6, -23.9]| 0|
| 3|29.5| -19.0| [29.5, -19.0]| 0|
| 4|32.8| -18.84| [32.8, -18.84]| 0|
+-----+-----+-----+-----+
```

Results and Discussions:

